

Yuming Liu PS6

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from pandas.plotting import scatter_matrix
import statsmodels.api as sm
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
```

Problem 1(a)

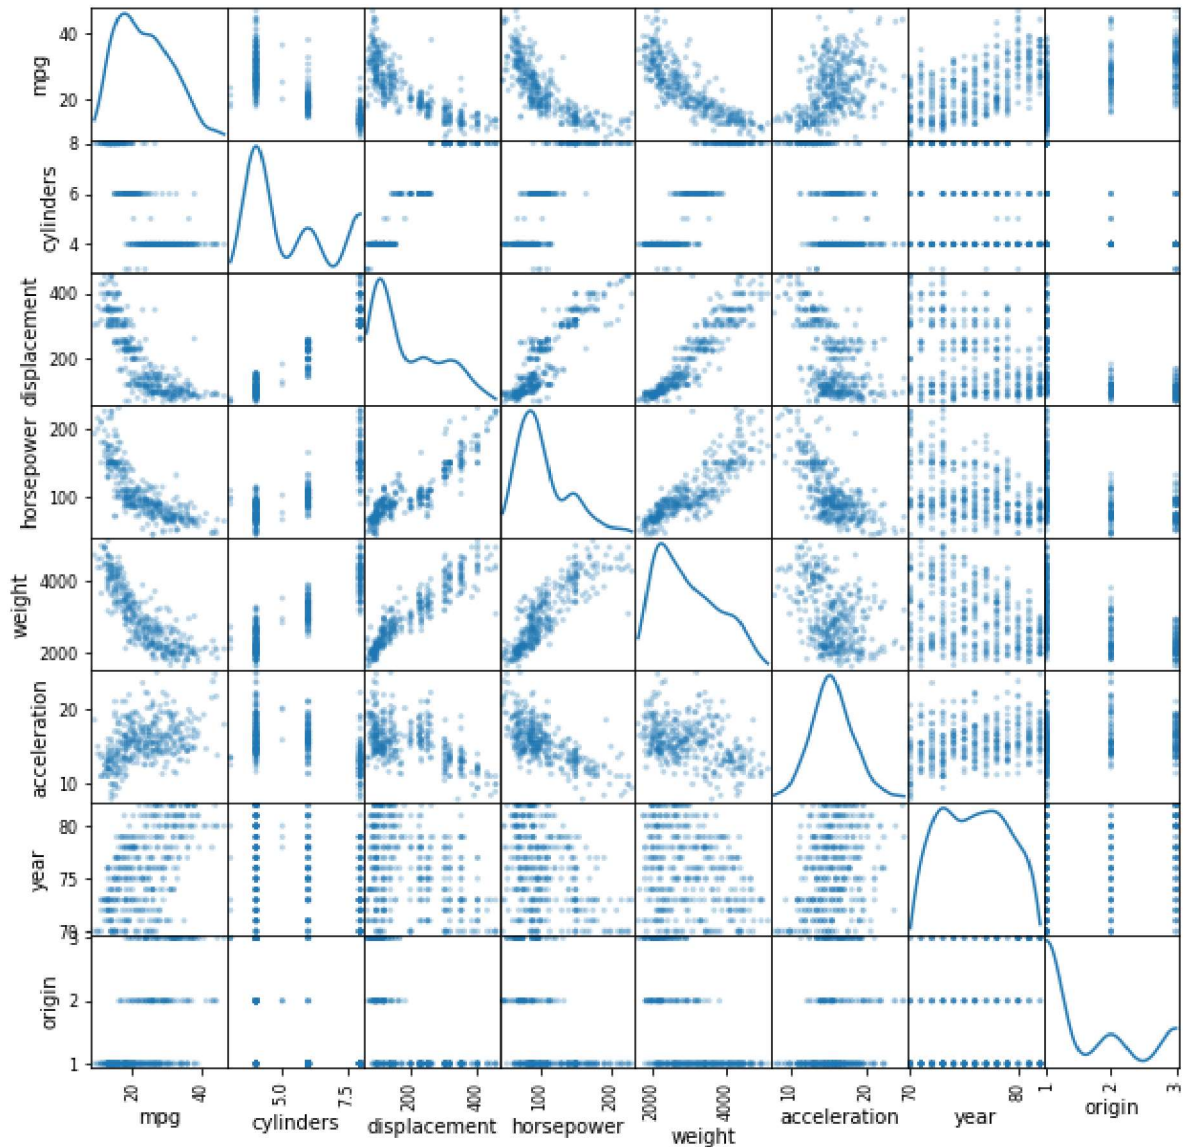
```
In [2]: df = pd.read_csv("data/Auto.csv", na_values='?')
df.dropna(inplace = True)
df.head()
```

Out[2]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name
0	18.0	8	307.0	130.0	3504	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150.0	3436	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150.0	3433	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140.0	3449	10.5	70	1	ford torino

Problem 1(b)

```
In [3]: df_quant = df[['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin']]
scatter_matrix(df_quant, alpha=0.3, ax=None, figsize=(10, 10), diagonal='kde')
plt.show()
```



Problem 1(c)

In [4]: `df_quant.corr()`

Out[4]:

	mpg	cylinders	displacement	horsepower	weight	acceleration	year
mpg	1.000000	-0.777618	-0.805127	-0.778427	-0.832244	0.423329	0.580541
cylinders	-0.777618	1.000000	0.950823	0.842983	0.897527	-0.504683	-0.345647
displacement	-0.805127	0.950823	1.000000	0.897257	0.932994	-0.543800	-0.369855
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	-0.689196	-0.416361
weight	-0.832244	0.897527	0.932994	0.864538	1.000000	-0.416839	-0.309120
acceleration	0.423329	-0.504683	-0.543800	-0.689196	-0.416839	1.000000	0.290316
year	0.580541	-0.345647	-0.369855	-0.416361	-0.309120	0.290316	1.000000
origin	0.565209	-0.568932	-0.614535	-0.455171	-0.585005	0.212746	0.181528

Problem 1(d)

In [5]: `df_quant['const'] = 1`

```
In [6]: reg1 = sm.OLS(endog=df_quant['mpg'], exog=df_quant[['const', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin']], missing='drop')
results = reg1.fit()
print(results.summary())
```

OLS Regression Results

```

=====
=
Dep. Variable:          mpg    R-squared:                0.82
1
Model:                  OLS    Adj. R-squared:           0.81
8
Method:                 Least Squares    F-statistic:           252.
4
Date:                   Sun, 16 Feb 2020    Prob (F-statistic):      2.04e-13
9
Time:                   17:18:40    Log-Likelihood:          -1023.
5
No. Observations:       392    AIC:                    206
3.
Df Residuals:           384    BIC:                    209
5.
Df Model:                7
Covariance Type:        nonrobust
=====

```

```

===
               coef      std err          t      P>|t|      [0.025      0.9
75]
-----
---
const          -17.2184      4.644      -3.707      0.000     -26.350     -8.
087
cylinders       -0.4934      0.323      -1.526      0.128     -1.129      0.
142
displacement     0.0199      0.008       2.647      0.008      0.005      0.
035
horsepower      -0.0170      0.014      -1.230      0.220     -0.044      0.
010
weight          -0.0065      0.001     -9.929      0.000     -0.008     -0.
005
acceleration     0.0806      0.099       0.815      0.415     -0.114      0.
275
year             0.7508      0.051     14.729      0.000      0.651      0.
851
origin           1.4261      0.278       5.127      0.000      0.879      1.
973
=====

```

```

=
Omnibus:           31.906    Durbin-Watson:           1.30
9
Prob(Omnibus):     0.000    Jarque-Bera (JB):         53.10
0
Skew:              0.529    Prob(JB):                 2.95e-1
2
Kurtosis:          4.460    Cond. No.                  8.59e+0
4
=====
=

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correc
tly specified.

```

[2] The condition number is large, $8.59e+04$. This might indicate that there are strong multicollinearity or other numerical problems.

- i. We have 'displacement', 'year', 'weight', and 'origin' are significant at 1% level.
- ii. We have 'horsepower', 'cylinders', and 'acceleration' are not significant at 10% level.
- iii. Suppose other variables will not change. We have 1 unit of year increase would bring mpg about 0.7508 unit increase.

Problem 1(e)

```
In [7]: df_quant['displacement2'] = np.square(df_quant['displacement'])
df_quant['horsepower2'] = np.square(df_quant['horsepower'])
df_quant['acceleration2'] = np.square(df_quant['acceleration'])
df_quant['weight2'] = np.square(df_quant['weight'])
```

```
In [8]: reg2 = sm.OLS(endog=df_quant['mpg'], exog=df_quant[['const', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin', 'displacement2', 'horsepower2', 'weight2', 'acceleration2']], missing='drop')
results2 = reg2.fit()
print(results2.summary())
```

OLS Regression Results

```

=====
=
Dep. Variable:          mpg    R-squared:                0.87
0
Model:                  OLS    Adj. R-squared:            0.86
6
Method:                 Least Squares    F-statistic:            230.
2
Date:                   Sun, 16 Feb 2020    Prob (F-statistic):      1.75e-16
0
Time:                   17:18:40    Log-Likelihood:          -962.0
2
No. Observations:       392    AIC:                    194
8.
Df Residuals:           380    BIC:                    199
6.
Df Model:                11
Covariance Type:        nonrobust
=====

```

```

=====
====
              coef      std err          t      P>|t|      [0.025      0.
975]
-----
----
const          20.1084      6.696       3.003     0.003      6.943      3
3.274
cylinders       0.2519      0.326       0.773     0.440     -0.389
0.893
displacement   -0.0169      0.020      -0.828     0.408     -0.057
0.023
horsepower     -0.1635      0.041     -3.971     0.000     -0.244     -
0.083
weight         -0.0136      0.003     -5.069     0.000     -0.019     -
0.008
acceleration   -2.0884      0.557     -3.752     0.000     -3.183     -
0.994
year           0.7810      0.045     17.512     0.000      0.693
0.869
origin         0.6104      0.263       2.320     0.021      0.093
1.128
displacement2  2.257e-05   3.61e-05     0.626     0.532   -4.83e-05   9.35
e-05
horsepower2    0.0004      0.000       2.943     0.003      0.000
0.001
weight2        1.514e-06   3.69e-07     4.105     0.000   7.89e-07   2.24
e-06
acceleration2  0.0576      0.016       3.496     0.001      0.025
0.090
=====
=
Omnibus:          33.614    Durbin-Watson:          1.57
6
Prob(Omnibus):    0.000    Jarque-Bera (JB):       77.98
5
Skew:             0.438    Prob(JB):               1.16e-1
7

```


Kurtosis: 5.002 Cond. No. 5.13e+08

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.13e+08. This might indicate that there are strong multicollinearity or other numerical problems.

ii. The adjusted R-squared stats is better than which of part(d).

iii. The terms are both non-significant at 10% level.

iv. It is not significant at 10% level, and its p-value is greater than which of the previous model.

Problem 1(f)

```
In [9]: X = [1, 6, 200, 100, 3100, 15.1, 99, 1, 200**2, 100**2, 3100**2, 15.1**2]
prediction = results2.predict(X)
print('The prediction mpg of cylinders displacement of 200, horsepower of 100,
a weight of 3,100, acceleration of 15.1, model year of 1999, and origin of 1 is
s', prediction[0])
```

The prediction mpg of cylinders displacement of 200, horsepower of 100, a weight of 3,100, acceleration of 15.1, model year of 1999, and origin of 1 is 38.73211109753366

Problem 2(a)

```
In [10]: df2 = pd.DataFrame({'X1':[0, 2, 0, 0, -1, 1], 'X2':[3, 0, 1, 1, 0, 1], 'X3':[0, 0, 3, 2, 1, 1], 'Y':['Red']*3+['Green']*2+['Red']})
```

```
In [11]: df2['Eucl. Dist from X1=X2=X3=0'] = np.sqrt(df2['X1']**2+df2['X2']**2+df2['X3']**2)
```

In [12]: df2

Out[12]:

	X1	X2	X3	Y	Eucl. Dist from X1=X2=X3=0
0	0	3	0	Red	3.000000
1	2	0	0	Red	2.000000
2	0	1	3	Red	3.162278
3	0	1	2	Green	2.236068
4	-1	0	1	Green	1.414214
5	1	1	1	Red	1.732051

Problem 2(b)

In [13]: `min(df2['Eucl. Dist from X1=X2=X3=0'])`

Out[13]: 1.4142135623730951

Since the fifth row has distance closest to $X1 = X2 = X3 = 0$, we have the prediction is green.

Problem 2(c)

Since the second, fifth, and sixth rows have distances closest to $X1 = X2 = X3 = 0$, we have the prediction is more likely to be red.

Problem 2(d)

If the Bayes (optimal) decision boundary in this problem is highly nonlinear, we would expect the best value for K to be large. Larger K can cover more points near the target, so the prediction could be more accurate.

Problem 2(e)

```
In [14]: neigh = KNeighborsClassifier(n_neighbors=2)
X = df2[['X1', 'X2', 'X3']]
Y = df2['Y']
neigh.fit(X, Y)
print('The KNN prediction for X1=X2=X3=1 and K=2 is', neigh.predict([(1,1,1)])
[0])
```

The KNN prediction for $X1=X2=X3=1$ and $K=2$ is Green

Problem 3(a)

```
In [15]: df_quant['mpg_high'] = np.where(df_quant['mpg'] >= np.median(df['mpg']), 1, 0)
```

```
In [16]: reg3 = sm.Logit(endog=df_quant['mpg_high'], exog=df_quant[['const', 'cylinders',
    'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin']],
    missing='drop')
    results3 = reg3.fit()
    print(results3.summary())
```

Optimization terminated successfully.

Current function value: 0.200944

Iterations 9

Logit Regression Results

```
=====
=
Dep. Variable:          mpg_high    No. Observations:          39
2
Model:                  Logit      Df Residuals:              38
4
Method:                 MLE        Df Model:
7
Date:                   Sun, 16 Feb 2020    Pseudo R-squ.:          0.710
1
Time:                   17:18:40    Log-Likelihood:         -78.77
0
converged:              True        LL-Null:                -271.7
1
Covariance Type:        nonrobust    LLR p-value:            2.531e-7
9
=====
===
              coef      std err          z      P>|z|      [0.025      0.9
75]
-----
---
const          -17.1549      5.764      -2.976      0.003     -28.452     -5.
858
cylinders       -0.1626      0.423      -0.384      0.701      -0.992      0.
667
displacement     0.0021      0.012       0.174      0.862      -0.021      0.
026
horsepower      -0.0410      0.024      -1.718      0.086      -0.088      0.
006
weight          -0.0043      0.001      -3.784      0.000      -0.007     -0.
002
acceleration     0.0161      0.141       0.114      0.910      -0.261      0.
293
year             0.4295      0.075       5.709      0.000       0.282      0.
577
origin           0.4773      0.362       1.319      0.187      -0.232      1.
187
=====
===
```

Possibly complete quasi-separation: A fraction 0.14 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

We have weight and year are significant at 5% level.

Problem 3(b)

```
In [17]: X = df_quant[['const','cylinders', 'displacement', 'horsepower', 'weight', 'ac
          celeration', 'year', 'origin']]
          Y = df_quant['mpg_high']

          X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.5, random_state=10)
```

Problem 3(c)

```
In [18]: clf = LogisticRegression(max_iter = 10000).fit(X_train, y_train)

          for i in range(7):
              print('The coefficient for',X.columns[i],'is',clf.coef_[0][i])
```

```
The coefficient for const is -0.0014062712185044482
The coefficient for cylinders is -1.1505902795202694
The coefficient for displacement is 0.01692195670266103
The coefficient for horsepower is 0.014535241600150541
The coefficient for weight is -0.007220539658275798
The coefficient for acceleration is 0.1521838862892103
The coefficient for year is 0.5780143927460908
```

Problem 3(d)

```
In [19]: predict_y = clf.predict(X_test)
          compare_y = confusion_matrix(y_test, predict_y)
```

```
In [20]: compare_y
```

```
Out[20]: array([[85, 14],
                [ 9, 88]], dtype=int64)
```

```
In [21]: print(classification_report(y_test, predict_y))
```

	precision	recall	f1-score	support
0	0.90	0.86	0.88	99
1	0.86	0.91	0.88	97
accuracy			0.88	196
macro avg	0.88	0.88	0.88	196
weighted avg	0.88	0.88	0.88	196

The model predicts both equally well.