

PS7 Yuming Liu

```
In [1]: import numpy as np
import pandas as pd
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, LeaveOneOut, KFold
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt
from scipy.interpolate import LSQUnivariateSpline

import warnings
warnings.filterwarnings("ignore")
```

Problem 1(a)

```
In [2]: df = pd.read_csv('data/strongdrink.txt')

X = df[['alco', 'malic', 'tot_phen', 'color_int']]
y = df['cultivar']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state=20)

result = LogisticRegression(solver='lbfgs', multi_class='multinomial').fit(X_train, y_train)

Results = pd.DataFrame({"j = 1": np.append(result.intercept_[0], result.coef_[0]),
                        "j = 2": np.append(result.intercept_[1], result.coef_[1])},
                        index=["beta_0", "beta_1", "beta_2", "beta_3", "beta_4"])
print(Results)

y_pred = result.predict(X_test)

print(classification_report(y_test, y_pred))
```

	j = 1	j = 2			
beta_0	-24.027617	22.780733			
beta_1	1.701734	-1.466297			
beta_2	-0.265788	-0.332951			
beta_3	1.224101	0.663556			
beta_4	0.022507	-0.922682			
	precision	recall	f1-score	support	
1	0.87	1.00	0.93	13	
2	1.00	0.90	0.95	21	
3	1.00	1.00	1.00	10	
accuracy			0.95	44	
macro avg	0.96	0.97	0.96	44	
weighted avg	0.96	0.95	0.96	44	

From the report, the error rates for the above groups are 13%, 0%, and 0%. Based on the f1-score, we have that $j = 3$ is the best model for prediction.

```
In [3]: MSE = (y_test != y_pred).mean()
print('The MSE from the test set is ', MSE)
```

The MSE from the test set is 0.045454545454545456

Problem 1(b)

```

In [4]: X = df[["alco", "malic", "tot_phen", "color_int"]].values
        y = df["cultivar"].values

        N_loo = X.shape[0]
        loo = LeaveOneOut()
        loo.get_n_splits(X)
        MSE_vec = np.zeros(N_loo)

        ypred = np.zeros(X.shape[0])

        for train_index, test_index in loo.split(X):
            X_train, X_test = X[train_index], X[test_index]
            y_train, y_test = y[train_index], y[test_index]
            LogReg = LogisticRegression(solver='lbfgs', multi_class='multinomial')
            LogReg.fit(X_train, y_train)
            y_pred = LogReg.predict(X_test)
            ypred[test_index] = y_pred
            if y_test == y_pred:
                MSE_vec[test_index] = 0
            else:
                MSE_vec[test_index] = 1

```

```

In [5]: print(classification_report(y, ypred))
        MSE_loo = MSE_vec.mean()
        print('The estimate MSE loocv of the test =', MSE_loo)

```

	precision	recall	f1-score	support
1	0.90	0.93	0.92	59
2	0.91	0.90	0.91	71
3	0.96	0.93	0.95	46
accuracy			0.92	176
macro avg	0.92	0.92	0.92	176
weighted avg	0.92	0.92	0.92	176

The estimate MSE loocv of the test = 0.07954545454545454

From the report, the error rates for the above groups are 10%, 9%, and 4%. The rates are higher than those from (a). Based on the f1-score, we have that $j = 3$ is the best model for prediction.

Problem 1(c)

```

In [6]: kf = KFold(n_splits=4, shuffle=True, random_state=10)
kf.get_n_splits(X)

MSE_vec_kf = np.zeros(4)

k_ind = int(0)
ypred = np.zeros(X.shape[0])

for train_index, test_index in kf.split(X):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
    LogReg = LogisticRegression()
    LogReg.fit(X_train, y_train)
    y_pred = LogReg.predict(X_test)
    ypred[test_index] = y_pred
    MSE_vec_kf[k_ind] = ((y_test - y_pred) ** 2).mean()
    k_ind += 1

```

```

In [7]: MSE_kf = MSE_vec_kf.mean()
print(classification_report(y, ypred))
print('The estimate MSE loocv of the test =', MSE_kf)

```

	precision	recall	f1-score	support
1	0.87	0.93	0.90	59
2	0.91	0.87	0.89	71
3	0.96	0.93	0.95	46
accuracy			0.91	176
macro avg	0.91	0.91	0.91	176
weighted avg	0.91	0.91	0.91	176


The estimate MSE loocv of the test = 0.10795454545454546

From the report, the error rates for the above groups are 13%, 9%, and 4%. The rates of 2 and 3 are higher than those from (a). The rate of 1 is higher than which from (b). The rates of 2 and 3 are possibly the same from (b). Based on the f1-score, we have that $j = 3$ is the best model for prediction.

Problem 2(a)

```
In [8]: df2 = pd.read_csv("data/CoolIndex.txt", names=["Age", "Cool"])

plt.scatter(x=df2['Age'], y=df2['Cool'], alpha=0.3, label="Points")
plt.title('Coolness Index data by age')
plt.legend()
plt.xlabel('Age')
plt.ylabel('Coolness Index')
plt.show()
```



Problem 2(b)

```
In [9]: df2["bin1"] = np.where((df2.Age >= 11) & (df2.Age < 22), 1, 0)
df2["bin2"] = np.where((df2.Age >= 22) & (df2.Age < 40), 1, 0)
df2["bin3"] = np.where((df2.Age >= 40) & (df2.Age < 59), 1, 0)
df2["bin4"] = np.where((df2.Age >= 59) & (df2.Age < 77), 1, 0)
df2["bin5"] = np.where((df2.Age >= 77) & (df2.Age <= 95), 1, 0)
```

```
In [10]: X = df2[["bin1", "bin2", "bin3", "bin4", "bin5"]]
y = df2['Cool']
res = sm.OLS(y, X, missing='drop').fit()
```


```
In [11]: print(res.summary())
```

```

                                OLS Regression Results
=====
=
Dep. Variable:                  Cool    R-squared:                  0.42
9
Model:                          OLS    Adj. R-squared:             0.42
7
Method:                          Least Squares    F-statistic:              178.
7
Date:                            Tue, 25 Feb 2020    Prob (F-statistic):       3.73e-11
4
Time:                            15:25:31    Log-Likelihood:           -3214.
5
No. Observations:                956    AIC:                      643
9.
Df Residuals:                    951    BIC:                      646
3.
Df Model:                        4
Covariance Type:                 nonrobust
=====
=
                                coef    std err          t      P>|t|      [0.025    0.97
5]
-----
-
bin1          20.1025      0.562     35.746     0.000     18.999     21.20
6
bin2          34.4758      0.431     80.006     0.000     33.630     35.32
1
bin3          37.6351      0.424     88.814     0.000     36.804     38.46
7
bin4          35.2254      0.485     72.560     0.000     34.273     36.17
8
bin5          27.2964      0.936     29.175     0.000     25.460     29.13
2
=====
=
Omnibus:                        80.102    Durbin-Watson:             1.23
6
Prob(Omnibus):                  0.000    Jarque-Bera (JB):          101.71
8
Skew:                           0.714    Prob(JB):                  8.17e-2
3
Kurtosis:                      3.719    Cond. No.                  2.2
1
=====
=

Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correc
tly specified.
```

```
In [12]: plt.scatter(df2['Age'], df2['Cool'], alpha=0.3, label='Points')
plt.plot(df2['Age'], res.predict(), 'r', label='OLS Step')
plt.legend()
plt.xlabel('Age')
plt.ylabel('Coolness Index')
plt.title('Coolness Index data by age')
plt.show()
```



```
In [13]: print('The predicted coolness of a 73-year old from the step function is', res
.params[3])
```


The predicted coolness of a 73-year old from the step function is 35.22540004
024275

Problem 2(c)

```
In [14]: df3 = df2.groupby('Age').mean()
df3['Age']=df3.index
```

```
In [15]: lsq = LSQUnivariateSpline(np.array(df3['Age']), np.array(df3['Cool']), t = [22
,40,59,77], k = 1)
```

```
In [16]: plt.scatter(df2['Age'], df2['Cool'], alpha=0.3, label='Points')
plt.plot(df2['Age'], res.predict(), 'r', label='OLS Step')
plt.plot(df3['Age'], lsq(df3['Age']), 'yellow', label='LSQ')
plt.legend()
plt.xlabel('Age')
plt.ylabel('Coolness Index')
plt.title('Coolness Index data by age')
plt.show()
```



```
In [17]: print('The predicted coolness of a 73-year old from the step function is', lsq(73))
```


The predicted coolness of a 73-year old from the step function is 32.86784862349653

Problem 2(d)

```
In [18]: lsq_new=LSQUnivariateSpline(np.array(df3['Age']), np.array(df3['Cool']), t=[22,40,59,77], k=3)
```



```
In [19]: plt.scatter(df2['Age'], df2['Cool'], alpha=0.3, label='Points')
plt.plot(df2['Age'], res.predict(), 'r', label='OLS Step')
plt.plot(df3['Age'], lsq(df3['Age']), 'yellow', label='LS')
plt.plot(df3['Age'], lsq_new(df3['Age']), 'green', label='Cubic Spline')
plt.legend()
plt.xlabel('Age')
plt.ylabel('Coolness Index')
plt.title('Coolness Index data by age')
plt.show()
```



```
In [20]: print('The predicted coolness of a 73-year old from the step function is', lsq
_new(73))
```

The predicted coolness of a 73-year old from the step function is 32.64230106
6279764